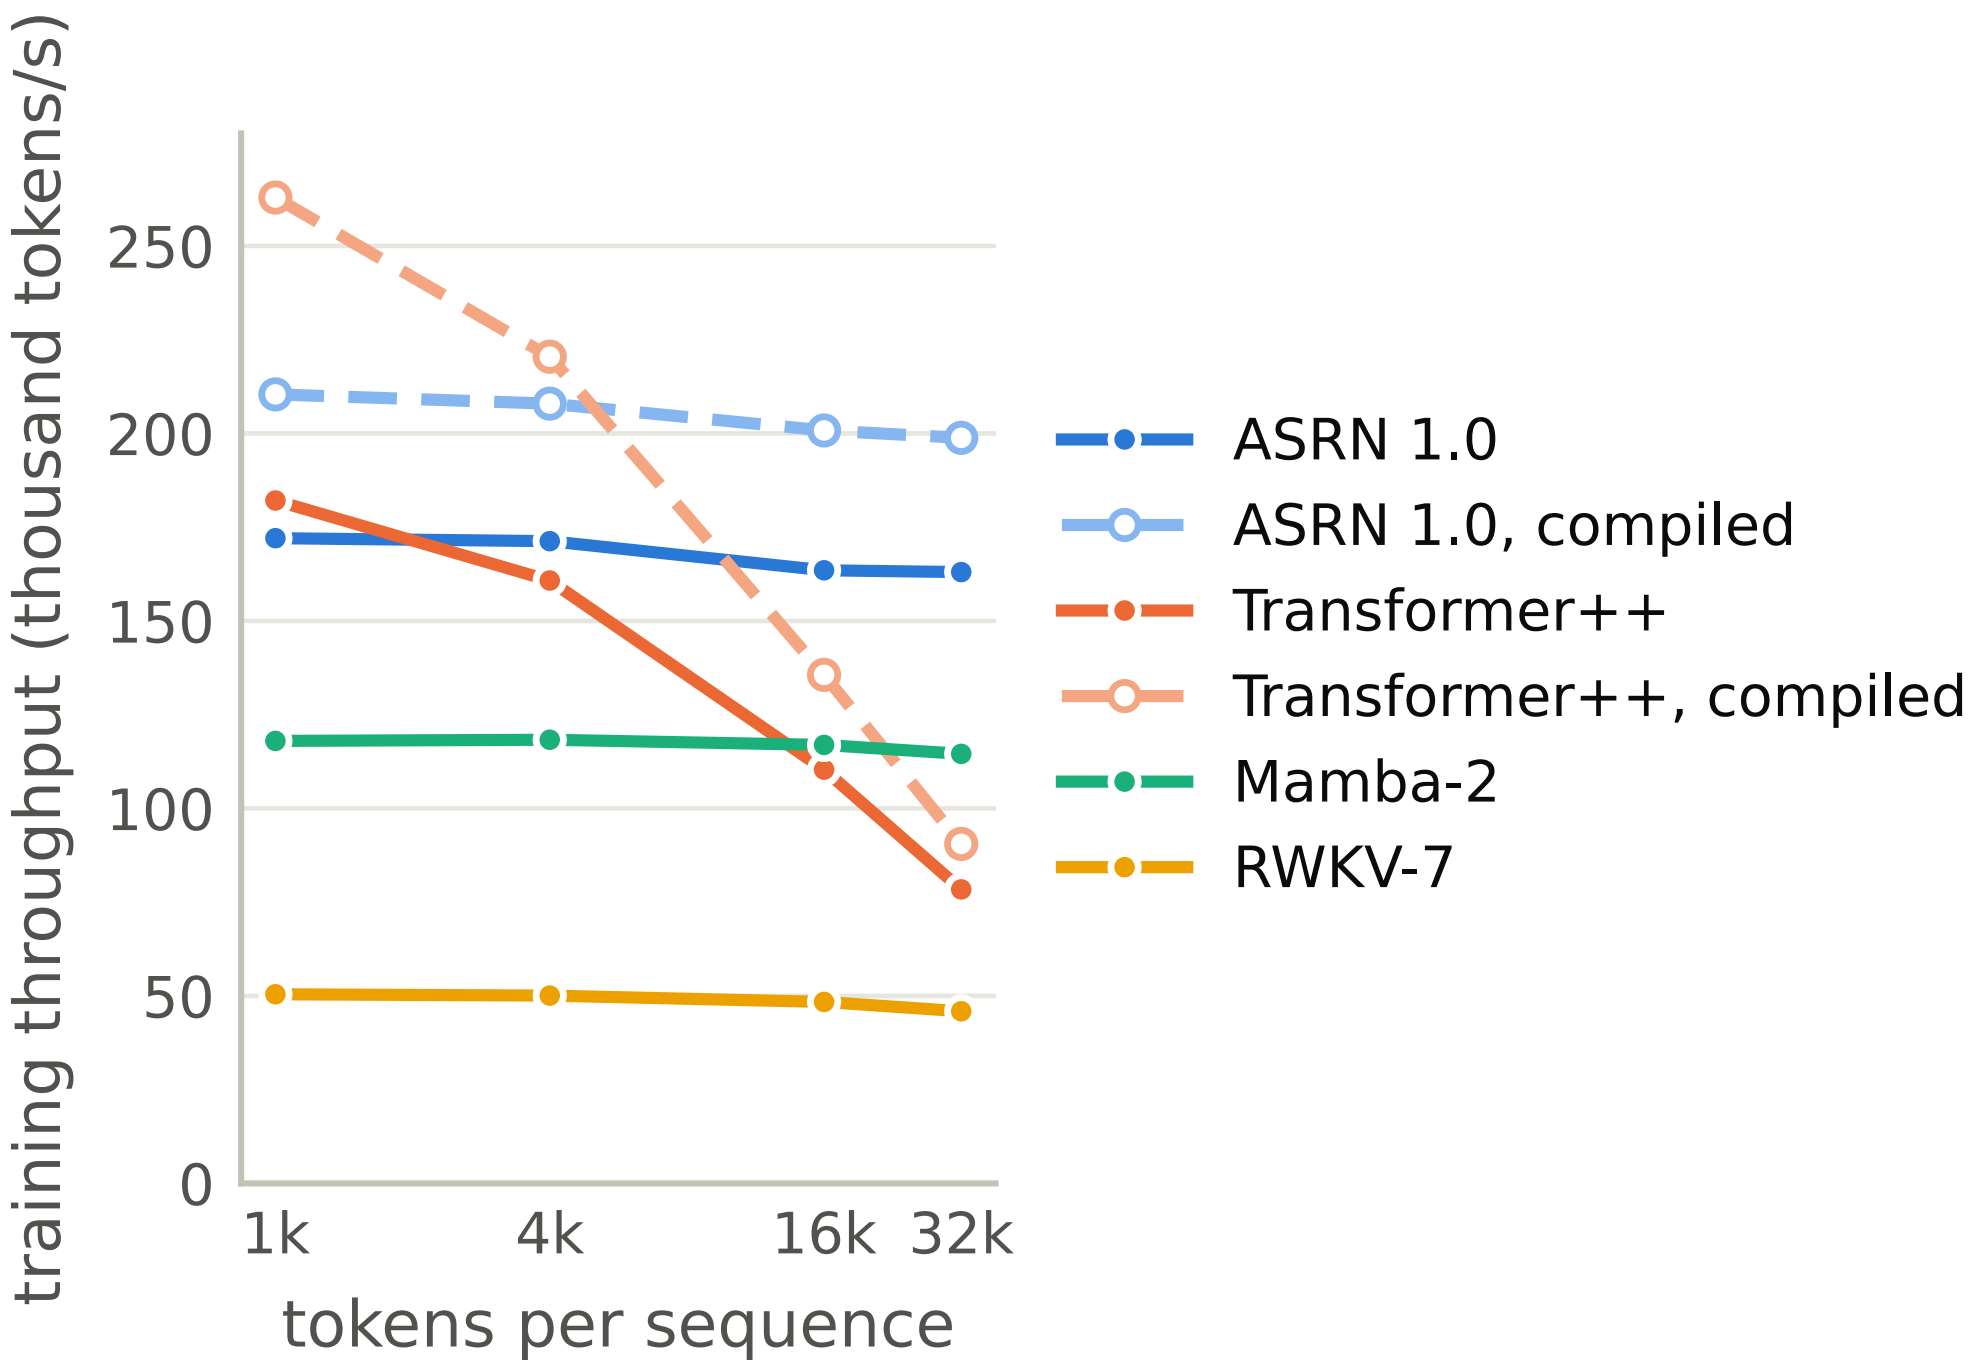
training throughput (thousand tokens/s)
250
200
150
100
50
0
1k
4k
16k
32k
tokens per sequence
ASRN 1.0
ASRN 1.0, compiled
Transformer++
Transformer++, compiled
Mamba-2
RWKV-7